

Sequence read quality

What it means

How it is determined

How it is represented in data files

The basics

- Sequence read quality is evaluated on a **base-by-base basis**
 - The quality of each base call is evaluated by a number
 - I.e. **quality scores**
- Whole reads, or sections of reads, can therefore be assessed on the quality properties of their bases
- E.g. each 'window' of N consecutive bases can easily have its properties determined, such as:
 - Average quality
 - Lowest quality base in the window, etc
- These properties can be used to decide where to 'trim' reads to remove poor quality segments
 - And which reads to discard entirely

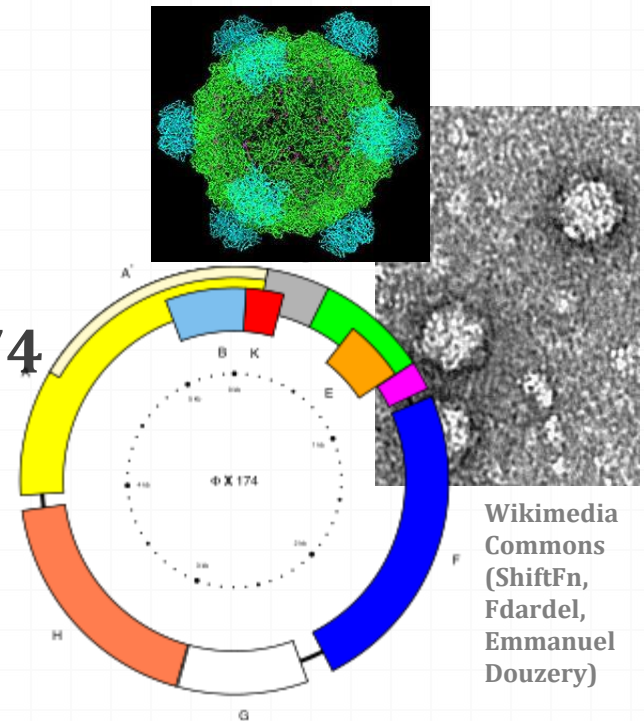


How sequence quality is determined



- A gold-standard reference sequence
 - i.e. you know the sequence for certain
- An objective way of measuring properties of your sequencing readout
- As an example, we will consider –
 - Good, old fashioned, **Sanger sequencing**
- Why Sanger?
 - It's as good as any as understanding the principles which link sequencing readout metrics with the reliability of the sequence
- Similar principles apply to NGS platforms

The reference sequence

- You need to be certain of the sequence.
- Therefore, it needs to be a piece of DNA which has been sequenced many, many times
 - so that each base is in no doubt
- So ideally it won't be huge (or tiny)
 - and it will be easily maintainable
- Step forward **bacteriophage Φ X174**
- It's genome is only about 5Kb




- Having a known sequence enables benchmarking
 - Basically, evaluate the characteristics of the sequencing readout
- Measurable properties of the readout
- what are these properties when a based is called:
 - Correctly
 - Incorrectly





Example readout: Sanger

- One 'lane' for each of A, C, G, T
- (due to labelled terminating dNTPs)
- Really old-fashioned – an actual lane of a gel
- Superseded by capillary sequencing




① Reaction mixture



- Primer and DNA template
- DNA polymerase
- ddNTPs with fluorochromes
- dNTPs (dATP, dCTP, dGTP, and dTTP)




③ Capillary gel electrophoresis separation of DNA fragments




④ Laser detection of fluorochromes and computational sequence analysis



J. Walshaw, GHTS, IFR




○ A rather nice section of sequencing readout




Principles of this **benchmarking process**:


- This nice, sharp peak in the T lane has measurable properties
- So too do the readouts in the A, C, G lanes at the same point (they are all flat)
- The base call is obvious
- Using the **known reference sequence**, all of the peaks with identical properties can be evaluated
- How often does a peak with these properties identify the correct base in the reference?
- How often is it wrong?
- **Probabilities of error for each base call can therefore be calculated**
 - (strictly speaking, estimated)



- What about this situation? A peak in two lanes simultaneously, albeit one is much smaller
- How often does a peak pair with these properties identify the correct base?
- How often is it wrong?
- What about these peaks:



- Similar principles can be applied to completely different sequencing platforms
- E.g. Illumina
- Again, it's a case of association readout characteristics with probabilities of error (error rates)



When is benchmarking done?

- In principle, it could be done as a one-off exercise
- In practice, it may make more sense for this to be done frequently
- i.e. as part of the **calibration** process prior to any sequencing run
- This is common in an Illumina sequencing operation
 - And uses sequences from the same Φ X genome
- So, need never concern you (unless you are actually operating the sequencing machine)

Using the benchmark error rates

- Once the error rates are known, these can then be applied to new output, i.e. of new, unknown sequences, base-call by base-call
 - I.e. one error probability per base
- The software which outputs the data from the sequencer does all this for you
- So again, no need to worry about **how** it's done
- But important to know what these probabilities **mean** – and how they are **expressed**

Phred scores

- Usually, the error probability (p) of each base call is expressed as a **Phred score**
- This value is simply: $-10 \times \log_{10} p$
 - rounded to an integer
- E.g. 1 in a thousand probability of being wrong
 - $p = 0.001$ (generally acceptable) \rightarrow phred score = 30
 - 1 in 10 (awful) : $p = 0.1$ \rightarrow phred score = 10
 - 1 in 4 (disastrous) : $p = 0.25$ \rightarrow phred score = 6
- Picking a base completely at random, i.e. 3 / 4 chance that it is wrong:
 - 3 in 4: $p = 0.75$ \rightarrow phred score = 1

Presenting the phred score for each base call

- QUAL file format: a PLAIN TEXT ('ASCII') file; each sequence has a header, then one digit per base

[illegible]

16/11/2016

J. Walshaw, GHFS, IFR

Presenting the phred score for each base call

o QUAL file format: header, then one digit per base

>IZFMVQQ01BF51O

40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40
40	38	38	38	40	40	40	39	39	39	34	34	34	40	40	40	40	40	40	40	40	40
39	40	40	40	40	40	40	40	40	39	26	26	26	40	40	39	39	39	40	40	40	40
40	40	40	40	40	40	40	40	39	21	21	21	35	40	40	40	40	40	40	40	40	40
33	33	40	40	40	40	40	40	40	40	40	40	40	40	34	34	34	40	40	40	35	27
40	40	40	40	40	40	40	40	40	39	39	39	40	40	40	40	40	40	40	40	40	40
22	19	24	24	36	36	36	36	36	32	33	34	35	40	40	39	39	30	28	28	40	33
22	18	18	18	31	36	36	36	36	36	36	36	36	31	31	31	31	31	31	34	34	31

>IZFMVQQ01A7E1H

40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40
40	40	39	39	39	40	40	40	21	21	21	39	39	39	40	40	40	40	40	40	40	40
40	40	40	40	40	40	40	40	40	40	40	40	40	40	21	21	21	30	39	39	39	39
40	40	40	40	40	40	40	40	39	39	39	40	40	40	39	39	39	40	40	40	40	40
38	39	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	39	39
40	40	40	40	40	40	40	40	38	38	38	40	40	40	40	40	40	40	40	40	40	40
39	39	40	40	40	40	40	34	34	35	40	40	40	40	40	40	40	40	40	40	34	34

16/11/2016

J. Walsh, GBS, QBR

QUAL and FASTA files

- Each QUAL file is used in conjunction with a FASTA file, which contains the corresponding base calls
 - I.e. the actual sequence
- The FASTA and QUAL files have corresponding headers

E.g. two sequence reads

>IZFMVQQ01BF510

40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40
40	38	38	38	40	40	40	39	39	39	34	34	34	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40
39	40	40	40	40	40	40	40	40	39	26	26	26	40	40	39	39	39	40	40	40	40	40	40	40	40	40	40	40
40	40	40	40	40	40	40	40	40	39	21	21	21	35	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40
33	33	40	40	40	40	40	40	40	40	40	40	40	40	34	34	34	40	40	40	35	27	27	27	30	35	40	40	31
40	40	40	40	40	40	40	40	40	39	39	39	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	35	28
22	19	24	24	36	36	36	36	36	32	33	34	35	40	40	39	39	30	28	28	40	33	33	33	33	35	13	26	26
22	18	18	18	31	36	36	36	36	36	36	36	36	31	31	31	31	31	31	31	31	31	31	31	36	27	31	31	31

>IZFMVQQ01A7E1H

40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40
40	40	39	39	39	40	40	40	21	21	21	39	39	39	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40
40	40	40	40	40	40	40	40	40	40	40	40	40	40	21	21	21	30	39	39	39	39	40	40	40	40	40	40	40	40
40	40	40	40	40	40	40	40	39	39	39	40	40	40	39	39	39	40	40	40	40	40	40	40	40	40	40	40	40	40
38	39	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	40	39									

QUAL

>IZFMVQQ01BF510

TCTCTATGCGGTGTCTAGCCGCCGCGGTAATACGTAGGGGCAAGCGTTATCCCGGATTTAC
TGGGTGTAAAGGGAGCGTAGACGGCAGCGCAAGTCTGAAGTGAAATGCCAGGGCTTAACC
CTGGAAGTGTCTTGGAACTGTGCAGCTAGAGTGCAGGAGAGGTAAGTGGAATTTCTAGT
GTAGCGGTGAAATGCGTAGATATTAGGAGGAACACCAGTGCGCGGAGGCGGCTTACTGGAC
GGTAACTGACGCTGAGGCTCGAAAGCGTGGGGAGCAAACAGGATTAGATACCCTGGTAGT
CCACGCCGTAAACGATGAATACTAGGTACAGGGGCACAAAAGTGCTTCTGTGCCGCAGCT
AACGCAATAAGTATTCCACCTGGGGAGTACGTTTCGCAAGAATGAAACTCAAAGGAATTGA
CGGGCTGAGACTGCCAAGGCACACAGGGGATAGGN

>IZFMVQQ01A7E1H

CGTGTCTCTAGTGCCAGCCGCCGCGGTAATACGTAGGTGGCAAGCGTTATCCCGGATTTA
CTGGGTGTAAAGGGCGTGTAGGCGGACGCTTAAGTCAGCGGTAAATTGCGGGGCTCAACC
TCGTCGAGCCGTTGAAACTGGGTGCCTTGAGTGGGCGAGAAGTACGCGGAATGCGTGGTG
TAGCGGTGAAATGCATAGATATCACGCAGAACTCCGATTGCGAAGGCAGCGTACCGGCGC
CCAAGTACGCTGAAGCACGAAGGCGTGGGTATCGAACAGGATTAGATACCTGGTAGTC
CACGCAGTAAACGATGAATGCTAGTTGTCCGGGGCGATTGAGTTCTGGGTGACACAGCGA
AAGCGTTAAGCATTCACCTGGGGAGTACGCCGGCAACGGTGAAACTTAAATGAATTGAC
GGGCTGAGACTGCCAAGGCACACAGGGGGATAGGN

FASTA

16/11/2016

FASTQ format

- A more commonly-used alternative to QUAL+FASTA
- Also a plain-text format
- Stores both the sequences and the quality (phred) scores in the same file
- Stores the quality scores in a more compact format than QUAL format
- Each possible Quality score is represented by a **single character** (letter, digit or symbol)
 - Thus, **encoding** of quality scores
 - As is sometimes the case with bioinformatics formats, things are not as simple as they might be:
 - There have been **different** encoding schemes employed

FASTQ format, and Example

○ 4 lines per read:

1. *@sequenceID additional-data*
2. *sequence (base calls) – one letter per base (of course...)*
3. *+sequenceID additional-data*
4. *encoded quality scores – one letter per base*

```
@HWI-M01242:112:000000000-AM193:1:1101:15596:1678 1:N:0:ACGCTACTGGATATCT
TACGGAGGATGCGAGCGTTATCCGATTATTGGGTTTAAAGGGTTCGTAGGCGGACGCTTAAGTCAG
+HWI-M01242:112:000000000-AM193:1:1101:15596:1678 1:N:0:ACGCTACTGGATATCT
BABBBFFDBFFFGGFGGGGAGGHFEAEBEFGHGFHGGGFDGGHGHGGGGGGGHGGEAGGGGGGGGHHHH
```

Note that lines 1 and 3 are necessarily identical (apart from the first character)

Example FASTQ

```
@HWI-M01242:112:000000000-AM193:1:1101:15596:1678 1:N:0:ACGCTACTGGATATCT
TACGGAGGATGCGAGCGTTATCCGGATTTATTGGGTTTAAAGGGTGCGTAGGCGGACGCTTAAGTCAGCGGTAAAATTGCGGGGCTCAACCTCC
+HWI-M01242:112:000000000-AM193:1:1101:15596:1678 1:N:0:ACGCTACTGGATATCT
BABBBFFDBFFFGGFGGGGAGGHFEAEBEFGHGFHGGGFDGGHGHGGGGGGGGHGGEGGGGGGGGHHHGGCEFFFGGGHFGGGGGGGGFCGHHG
@HWI-M01242:112:000000000-AM193:1:1101:16562:1707 1:N:0:ACGCTACTGGATATCT
TACGGAGGATACGAGCGTTATCCGGATTTATTGGGTTTAAAGGGTGCGTAGGCGGAGTGTCAAGTCAGCGGTAAAATTTTCGGGGCTCAACCCCC
+HWI-M01242:112:000000000-AM193:1:1101:16562:1707 1:N:0:ACGCTACTGGATATCT
BBBBBBBBBAFFFFGGGGGGEGGHGGAEGHGGHHHHEGGHHHHHGGGGGFGGGHGGFGGHHHHHHHHHHHGGGGGHHGHGHFGGGFGEHHHEGG
(...etc)
```

- (2 reads are shown above)

Example FASTQ

- In practice, the sequenceID and additional data is often omitted from line 3 to save space, with the 4th line being assumed to be the Q-scores of the most recent sequence
- (same 2 reads are shown below, with implicit headers prior to the Q-scores)

```
@HWI-M01242:112:000000000-AM193:1:1101:15596:1678 1:N:0:ACGCTACTGGATATCT
TACGGAGGATGCGAGCGTTATCCGGATTTATTGGGTTTAAAGGGTGCGTAGGCGGACGCTTAAGTCAGCGGTAAAATTGCGGGGCTCAACCTCC
+
BABBBFFDBFFFGGFGGGGAGGHFEAEBEFGHGFHGGGFDGGHGHGGGGGGGGHGGEGGGGGGGGHHHHGGCEFFFGGGGHFGGGGGGGGFCGHHG
@HWI-M01242:112:000000000-AM193:1:1101:16562:1707 1:N:0:ACGCTACTGGATATCT
TACGGAGGATACGAGCGTTATCCGGATTTATTGGGTTTAAAGGGTGCGTAGGCGGAGTGTCAAGTCAGCGGTAAAATTTTCGGGGCTCAACCCCC
+
BBBBBBBBAFFFFGGGGGGEGGHGGAEGHGGHHHHEGGHHHHHGGGGGFGGGHGGFGGHHHHHHHHHHHGGGGGHGHGHFGGGFGEHHHEG
```

Encoding quality-scores in FASTQ

- BABBBFFDBFFFGGFGGGGAGGGHFEAEBEFGHGFHGGGFDGGHG...
- So what do those quality-score lines mean?
 - Depending on the coding scheme, quality scores of up to **93** can be stored
 - In a read output by a sequencer, even the best scores you see in practice will be **far lower**
 - E.g. best scores of around **40** ('I' in the current Illumina format)
- For more details on encoding of quality scores and their relation to **ASCII codes**, refer to:
 - MF's slides ("*It's all about the text*") of 19th Oct:
http://ghfs1.ifr.ac.uk/ghfs/wp-content/uploads/2016/10/Bitesize_ngs_formats.pdf
 - JW's background ("*A brief guide to how computers encode data*"):
http://ghfs1.ifr.ac.uk/ghfs/wp-content/uploads/2016/11/slides_file_storage_1.pdf

Format conversion

- Sequencing providers give Illumina-format data in the form of FASTQ
- Many tools will understand FASTQ
- So conversion may not be necessary
- Some conversion tools can convert between FASTA+QUAL and FASTQ
 - E.g. PRINSEQ (prinseq-lite.pl)
- EMBOSS seqret (a very versatile program)
 - FASTQ→FASTA (loses quality scores)

Other formats you may have come across

- SFF : Standard Flowgram File
 - The native format of the 454 sequencing platform
 - Comes in binary and plain-text varieties
 - Can be converted to other formats such as FASTA/QUAL with various tools, e.g. MOTHUR sffinfo
- BCL: earlier Illumina Base Call format
 - With current Illumina software, BCL files will have been converted to FASTQ before you are likely to have seen the data
- FAST5: used with Oxford Nanopore sequencers
 - Based on HDF5 (Hierarchical Data Format – a generic and versatile data format definition)